

Sqoop Assignment

Task 1

Use Sqoop tool to export data present in SGOOPOUT folder made while demo of Import table

Before exporting data to hdfs using Sqoop, I have inserted data into Person table as below

```
Applications Places System
File Edit View Search Terminal Help
mysql> select * from Person;
Empty set (0.00 sec)

mysql>
mysql> insert into Person values (1,'Abhishek','Ram','Patna','Bihar');
Query OK, 1 row affected (0.01 sec)

mysql> insert into Person values (2,'Tanya','Priya','Whitefiled','Bangalore');
Query OK, 1 row affected (0.01 sec)

mysql> insert into Person values (3,'James','Brown','New York','United States');
Query OK, 1 row affected (0.02 sec)

mysql> insert into Person values (4,'Jhon','Miller','Los Angeles','United States');
Query OK, 1 row affected (0.00 sec)

mysql> select * from Person;
+-----+-----+-----+-----+-----+
| person_id | lname | fname | area | city |
+-----+-----+-----+-----+-----+
| 1 | Abhishek | Ram | Patna | Bihar |
| 2 | Tanya | Priya | Whitefiled | Bangalore |
| 3 | James | Brown | New York | United States |
| 4 | Jhon | Miller | Los Angeles | United States |
+-----+-----+-----+-----+-----+
4 rows in set (0.00 sec)
```

I have exported data to hdfs dir /Sqoop1 as below

```
Applications Places System
acacgild@192:~
File Edit View Search Terminal Help
States
[acacgild@192 ~]$ sqoop import --connect jdbc:mysql://localhost/simplydb --table Person --username root -P --target-dir /Sqoop1 -m 2 --split-by Person_Id
Warning: /home/acacgild/install/sqoop/sqoop-1.4.6/bin_hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT HOME to the root of your HCatalog installation.
Warning: /home/acacgild/install/sqoop/sqoop-1.4.6/bin_hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/06/06 00:25:47 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/06/06 00:25:53 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/06/06 00:25:53 INFO tool.CodeGenTool: Beginning code generation
Wed Jun 06 00:25:53 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/06/06 00:25:54 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person' AS t LIMIT 1
18/06/06 00:25:54 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person' AS t LIMIT 1
18/06/06 00:25:54 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acacgild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acacgild/compile/3e0185d0df900b5214887472d68edbf/Person.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/06/06 00:25:56 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acacgild/compile/3e0185d0df900b5214887472d68edbf/Person.jar
18/06/06 00:25:56 WARN manager.MySQLManager: It looks like you are importing from mysql.
18/06/06 00:25:56 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
18/06/06 00:25:56 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
18/06/06 00:25:56 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
18/06/06 00:25:56 INFO mapreduce.ImportJobBase: Beginning import of Person
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acacgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acacgild/install/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/06/06 00:25:56 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/06/06 00:25:56 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
18/06/06 00:25:57 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
18/06/06 00:25:57 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
Wed Jun 06 00:26:00 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/06/06 00:26:00 INFO db.DBInputFormat: Using read committed transaction isolation
18/06/06 00:26:00 INFO db.DataDrivenDBInputFormat: BoundingValsQuery: SELECT MIN('Person_Id'), MAX('Person_Id') FROM 'Person'
18/06/06 00:26:00 INFO mapreduce.JobSubmitter: number of splits:2
18/06/06 00:26:00 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1528221645710_0003
18/06/06 00:26:01 INFO impl.YarnClientImpl: Submitted application application_1528221645710_0003
18/06/06 00:26:01 INFO mapreduce.Job: The url to track the job: http://192.168.79.128:8088/proxy/application_1528221645710_0003/
18/06/06 00:26:01 INFO mapreduce.Job: Running job: job_1528221645710_0003
18/06/06 00:26:10 INFO mapreduce.Job: Job job_1528221645710_0003 running in uber mode : false
18/06/06 00:26:10 INFO mapreduce.Job: map 0% reduce 0%
18/06/06 00:26:21 INFO mapreduce.Job: map 50% reduce 0%
18/06/06 00:26:22 INFO mapreduce.Job: map 100% reduce 0%
18/06/06 00:26:22 INFO mapreduce.Job: Job job_1528221645710_0003 completed successfully
18/06/06 00:26:23 INFO mapreduce.Job: Counters: 30
File System Counters
FILE: Number of bytes read=0
FILE: Number of bytes written=254406
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
```

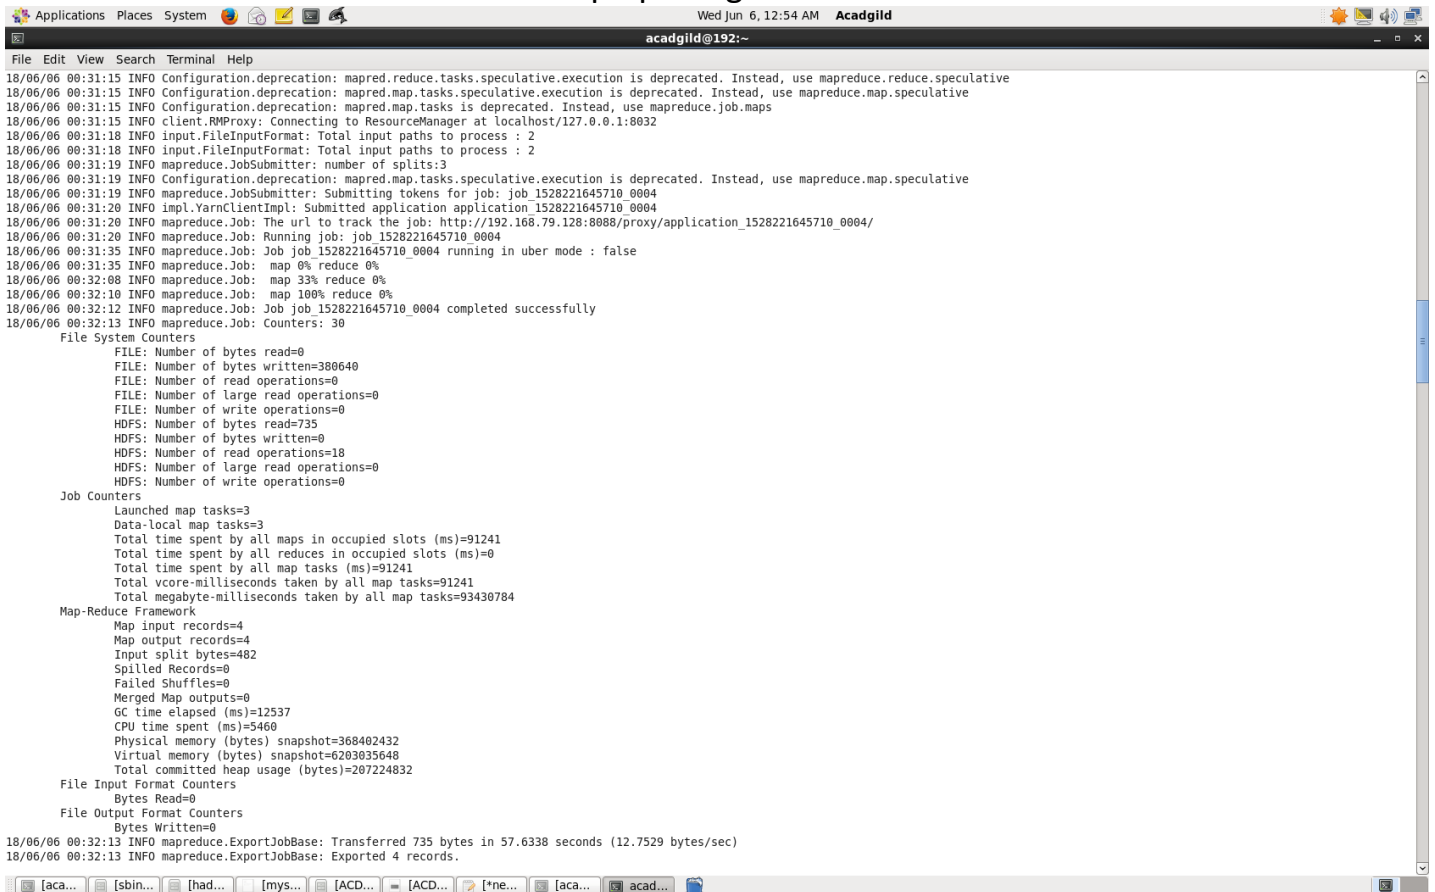
Sqoop Assignment

```
Applications Places System Wed Jun 6, 12:54 AM Acadgild
acadgild@192:~
File Edit View Search Terminal Help
18/06/06 00:26:21 INFO mapreduce.Job: map 50% reduce 0%
18/06/06 00:26:22 INFO mapreduce.Job: map 100% reduce 0%
18/06/06 00:26:22 INFO mapreduce.Job: Job job_1528221645710_0003 completed successfully
18/06/06 00:26:23 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=254406
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=225
    HDFS: Number of bytes written=141
    HDFS: Number of read operations=8
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=4
  Job Counters
    Launched map tasks=2
    Other Local map tasks=2
    Total time spent by all maps in occupied slots (ms)=18452
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=18452
    Total vcore-milliseconds taken by all map tasks=18452
    Total megabyte-milliseconds taken by all map tasks=18894848
  Map-Reduce Framework
    Map input records=4
    Map output records=4
    Input split bytes=225
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=499
    CPU time spent (ms)=2780
    Physical memory (bytes) snapshot=245317632
    Virtual memory (bytes) snapshot=4135526400
    Total committed heap usage (bytes)=138149888
  File Input Format Counters
    Bytes Read=0
  File Output Format Counters
    Bytes Written=141
18/06/06 00:26:23 INFO mapreduce.ImportJobBase: Transferred 141 bytes in 25.5998 seconds (5.5079 bytes/sec)
18/06/06 00:26:23 INFO mapreduce.ImportJobBase: Retrieved 4 records.
[acadgild@192 ~]$ hadoop fs -ls /Sqoop1
18/06/06 00:26:36 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
-rw-r--r-- 1 acadgild supergroup          0 2018-06-06 00:26 /Sqoop1/_SUCCESS
-rw-r--r-- 1 acadgild supergroup        62 2018-06-06 00:26 /Sqoop1/part-m-00000
-rw-r--r-- 1 acadgild supergroup        79 2018-06-06 00:26 /Sqoop1/part-m-00001
[acadgild@192 ~]$ hadoop fs -cat /Sqoop1/part-m-00000
18/06/06 00:26:41 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
^[[A][A1,Abhishek,Ram,Patna,Bihar
2,Tanya,Priya,Whitefiled,Bangalore
[acadgild@192 ~]$ hadoop fs -cat /Sqoop1/part-m-00001
18/06/06 00:26:49 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
3,James,Brown,New York,United States
4,Jhon,Miller, Los Angeles ,United States
```

Exported back from HDFS Dir /Sqoop1 to SqoopExport database where Person is a table in it.

```
Applications Places System Wed Jun 6, 12:54 AM Acadgild
acadgild@192:~
File Edit View Search Terminal Help
[acadgild@192 ~]$ sqoop export --connect jdbc:mysql://localhost/SqoopExport --username root -P -export-dir /Sqoop1
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/06/06 00:31:06 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/06/06 00:31:11 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/06/06 00:31:11 INFO tool.CodeGenTool: Beginning code generation
Wed Jun 06 00:31:11 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/06/06 00:31:12 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person' AS t LIMIT 1
18/06/06 00:31:12 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person' AS t LIMIT 1
18/06/06 00:31:12 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadgild/compile/40278ad73cd6d8b4bbe1715ec465d38/Person.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/06/06 00:31:14 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/40278ad73cd6d8b4bbe1715ec465d38/Person.jar
18/06/06 00:31:14 INFO mapreduce.ExportJobBase: Beginning export of Person
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/06/06 00:31:14 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/06/06 00:31:14 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
18/06/06 00:31:15 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.reduce.speculative
18/06/06 00:31:15 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
18/06/06 00:31:15 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
18/06/06 00:31:15 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/06/06 00:31:18 INFO input.FileInputFormat: Total input paths to process : 2
18/06/06 00:31:18 INFO input.FileInputFormat: Total input paths to process : 2
18/06/06 00:31:19 INFO mapreduce.JobSubmitter: number of splits:3
18/06/06 00:31:19 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
18/06/06 00:31:19 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1528221645710_0004
18/06/06 00:31:20 INFO Impl.YarnClientImpl: Submitted application application_1528221645710_0004
18/06/06 00:31:20 INFO mapreduce.Job: The url to track the job: http://192.168.79.128:8088/proxy/application_1528221645710_0004/
18/06/06 00:31:20 INFO mapreduce.Job: Running job: job_1528221645710_0004
18/06/06 00:31:35 INFO mapreduce.Job: Job job_1528221645710_0004 running in uber mode : false
18/06/06 00:31:35 INFO mapreduce.Job: map 0% reduce 0%
18/06/06 00:32:08 INFO mapreduce.Job: map 33% reduce 0%
18/06/06 00:32:10 INFO mapreduce.Job: map 100% reduce 0%
18/06/06 00:32:12 INFO mapreduce.Job: Job job_1528221645710_0004 completed successfully
18/06/06 00:32:13 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=380640
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=735
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=18
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=0
```

Sqoop Assignment



The screenshot shows a terminal window titled 'acadgild@192:~' with a menu bar (File, Edit, View, Search, Terminal, Help). The terminal displays the output of a Sqoop export command. The logs show the configuration of the export job, the submission of the application to the YARN cluster, and the progress of the map tasks. The job is identified as 'job_1528221645710_0004'. The logs indicate that the export was successful, with 4 records transferred from the HDFS directory 'sqoopexport' to the MySQL database table 'Person'.

```
18/06/06 00:31:15 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.reduce.speculative
18/06/06 00:31:15 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
18/06/06 00:31:15 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
18/06/06 00:31:15 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/06/06 00:31:18 INFO input.FileInputFormat: Total input paths to process : 2
18/06/06 00:31:18 INFO input.FileInputFormat: Total input paths to process : 2
18/06/06 00:31:19 INFO mapreduce.JobSubmitter: number of splits:3
18/06/06 00:31:19 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
18/06/06 00:31:19 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1528221645710_0004
18/06/06 00:31:20 INFO impl.YarnClientImpl: Submitted application application_1528221645710_0004
18/06/06 00:31:20 INFO mapreduce.Job: The url to track the job: http://192.168.79.128:8088/proxy/application_1528221645710_0004/
18/06/06 00:31:20 INFO mapreduce.Job: Running job: job_1528221645710_0004
18/06/06 00:31:35 INFO mapreduce.Job: Job job_1528221645710_0004 running in uber mode : false
18/06/06 00:31:35 INFO mapreduce.Job: map 0% reduce 0%
18/06/06 00:32:08 INFO mapreduce.Job: map 33% reduce 0%
18/06/06 00:32:10 INFO mapreduce.Job: map 100% reduce 0%
18/06/06 00:32:12 INFO mapreduce.Job: Job job_1528221645710_0004 completed successfully
18/06/06 00:32:13 INFO mapreduce.Job: Counters: 30
File System Counters
  FILE: Number of bytes read=0
  FILE: Number of bytes written=380640
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=735
  HDFS: Number of bytes written=0
  HDFS: Number of read operations=18
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=0
Job Counters
  Launched map tasks=3
  Data-local map tasks=3
  Total time spent by all maps in occupied slots (ms)=91241
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=91241
  Total vcore-milliseconds taken by all map tasks=91241
  Total megabyte-milliseconds taken by all map tasks=93430784
Map-Reduce Framework
  Map input records=4
  Map output records=4
  Input split bytes=482
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=12537
  CPU time spent (ms)=5460
  Physical memory (bytes) snapshot=368402432
  Virtual memory (bytes) snapshot=6203035648
  Total committed heap usage (bytes)=207224832
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=0
18/06/06 00:32:13 INFO mapreduce.ExportJobBase: Transferred 735 bytes in 57.6338 seconds (12.7529 bytes/sec)
18/06/06 00:32:13 INFO mapreduce.ExportJobBase: Exported 4 records.
```

Select * from Person table; SqoopExport database after loading data from HDFS Dir using Sqoop

```
mysql> use SqoopExport;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> select * from Person;
Empty set (0.00 sec)

mysql> select * from Person;
+-----+-----+-----+-----+-----+
| person_id | lname | fname | area | city |
+-----+-----+-----+-----+-----+
| 1 | Abhishek | Ram | Patna | Bihar |
| 2 | Tanya | Priya | Whitefield | Bangalore |
| 3 | James | Brown | New York | United States |
| 4 | Jhon | Miller | Los Angeles | United States |
+-----+-----+-----+-----+-----+
4 rows in set (0.00 sec)
```

Task2

Use Sqoop tool to export data present in SQOOPOUT folder made while demo of Import table with parameter person_id =3.

Exported data from database simplidb where personid=3 , data was already loaded earlier so, using the same

For this query.

... Continue below

Sqoop Assignment

```
Applications Places System [Icons] [Taskbar] [System Tray] [Network] [Volume] [Battery] [Clock]
acadgild@192:~
File Edit View Search Terminal Help
[acadgild@192 ~]$ sqoop import --connect jdbc:mysql://localhost/simplidb --table Person --username root -P --target-dir /Sqoop2 -m 1 --where Person_Id=3
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/06/06 00:39:24 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/06/06 00:39:28 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/06/06 00:39:28 INFO tool.CodeGenTool: Beginning code generation
Wed Jun 06 00:39:28 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be es
tablished by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by
setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/06/06 00:39:29 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person' AS t LIMIT 1
18/06/06 00:39:29 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person' AS t LIMIT 1
18/06/06 00:39:29 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadgild/compile/b4bc4a2f2504b616b7f8f681a3b37f8e/Person.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/06/06 00:39:32 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/b4bc4a2f2504b616b7f8f681a3b37f8e/Person.jar
18/06/06 00:39:32 WARN manager.MySQLManager: It looks like you are importing from mysql.
18/06/06 00:39:32 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
18/06/06 00:39:32 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
18/06/06 00:39:32 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
18/06/06 00:39:32 INFO mapreduce.ImportJobBase: Beginning import of Person
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/06/06 00:39:32 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/06/06 00:39:32 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
18/06/06 00:39:34 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
18/06/06 00:39:34 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
Wed Jun 06 00:39:41 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be es
tablished by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by
setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/06/06 00:39:41 INFO db.InputFormat: Using read committed transaction isolation
18/06/06 00:39:41 INFO mapreduce.JobSubmitter: number of splits:1
18/06/06 00:39:42 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1528221645710_0006
18/06/06 00:39:45 INFO impl.YarnClientImpl: Submitted application application_1528221645710_0006
18/06/06 00:39:45 INFO mapreduce.Job: The url to track the job: http://192.168.79.128:8088/proxy/application_1528221645710_0006/
18/06/06 00:39:45 INFO mapreduce.Job: Running job: job_1528221645710_0006
18/06/06 00:40:57 INFO mapreduce.Job: Job job_1528221645710_0006 running in uber mode : false
18/06/06 00:40:57 INFO mapreduce.Job: map 0% reduce 0%
18/06/06 00:40:58 INFO mapreduce.Job: map 100% reduce 0%
18/06/06 00:41:00 INFO mapreduce.Job: Job job_1528221645710_0006 completed successfully
18/06/06 00:41:00 INFO mapreduce.Job: Counters: 30
File System Counters
FILE: Number of bytes read=0
FILE: Number of bytes written=127351
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=87
HDFS: Number of bytes written=37
HDFS: Number of read operations=4
HDFS: Number of write operations=2
Job Counters
Launched map tasks=1
Other local map tasks=1
Total time spent by all maps in occupied slots (ms)=20792
Total time spent by all reduces in occupied slots (ms)=0
Total time spent by all map tasks (ms)=20792
Total vcore-milliseconds taken by all map tasks=20792
Total megabyte-milliseconds taken by all map tasks=21291008
Map-Reduce Framework
Map input records=1
Map output records=1
Input split bytes=87
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=88
CPU time spent (ms)=2130
Physical memory (bytes) snapshot=122527744
Virtual memory (bytes) snapshot=2063437824
Total committed heap usage (bytes)=69074944
File Input Format Counters
Bytes Read=0
File Output Format Counters
Bytes Written=37
18/06/06 00:41:00 INFO mapreduce.ImportJobBase: Transferred 37 bytes in 86.2087 seconds (0.4292 bytes/sec)
18/06/06 00:41:00 INFO mapreduce.ImportJobBase: Retrieved 1 records.
[acadgild@192 ~]$ hadoop fs -ls /Sqoop2
18/06/06 00:41:44 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup 0 2018-06-06 00:40 /Sqoop2/ SUCCESS
-rw-r--r-- 1 acadgild supergroup 37 2018-06-06 00:40 /Sqoop2/part-m-00000
[acadgild@192 ~]$ hadoop fs -cat /Sqoop2/part-m-00000
18/06/06 00:41:57 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
3,James,Brown,New York,United States
```


Sqoop Assignment

Exported data from HDFS Dir to SqoopExport2 data base & Person table, where person_id=3

```
Applications Places System Wed Jun 6, 12:55 AM Acadgild
acadmild@192:~
File Edit View Search Terminal Help
3,James,Brown,New York,United States
[acadmild@192 ~]$ sqoop export --connect jdbc:mysql://localhost/SqoopExport2 --table Person --username root -P --export-dir /Sqoop2
Warning: /home/acadmild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadmild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/06/06 00:44:28 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/06/06 00:44:33 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/06/06 00:44:33 INFO tool.CodeGenTool: Beginning code generation
Wed Jun 06 00:44:33 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be es
tablished by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by
setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/06/06 00:44:35 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person` AS t LIMIT 1
18/06/06 00:44:36 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person` AS t LIMIT 1
18/06/06 00:44:36 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadmild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadmild/compile/14ef46086689dc29adfe89fe2dd90bd0/Person.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/06/06 00:44:39 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadmild/compile/14ef46086689dc29adfe89fe2dd90bd0/Person.jar
18/06/06 00:44:39 INFO mapreduce.ExportJobBase: Beginning export of Person
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadmild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadmild/install/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/06/06 00:44:40 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/06/06 00:44:40 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
18/06/06 00:44:42 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.reduce.speculative
18/06/06 00:44:42 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
18/06/06 00:44:42 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
18/06/06 00:44:42 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/06/06 00:44:45 INFO input.FileInputFormat: Total input paths to process : 1
18/06/06 00:44:45 INFO input.FileInputFormat: Total input paths to process : 1
18/06/06 00:44:45 INFO mapreduce.JobSubmitter: number of splits:4
18/06/06 00:44:45 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
18/06/06 00:44:46 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1528221645710_0007
18/06/06 00:44:46 INFO impl.YarnClientImpl: Submitted application application_1528221645710_0007
18/06/06 00:44:46 INFO mapreduce.Job: The url to track the job: http://192.168.79.128:8088/proxy/application_1528221645710_0007/
18/06/06 00:44:46 INFO mapreduce.Job: Running job: job_1528221645710_0007
18/06/06 00:44:58 INFO mapreduce.Job: Job job_1528221645710_0007 running in uber mode : false
18/06/06 00:44:58 INFO mapreduce.Job: map 0% reduce 0%
18/06/06 00:45:40 INFO mapreduce.Job: map 25% reduce 0%
18/06/06 00:45:41 INFO mapreduce.Job: map 100% reduce 0%
18/06/06 00:45:41 INFO mapreduce.Job: Job job_1528221645710_0007 completed successfully
18/06/06 00:45:41 INFO mapreduce.Job: Counters: 30
File System Counters
FILE: Number of bytes read=0
FILE: Number of bytes written=507524
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=660
HDFS: Number of bytes written=0
HDFS: Number of read operations=19
HDFS: Number of large read operations=0

Job Counters
Launched map tasks=4
Data-local map tasks=4
Total time spent by all maps in occupied slots (ms)=160527
Total time spent by all reduces in occupied slots (ms)=0
Total time spent by all map tasks (ms)=160527
Total vcore-milliseconds taken by all map tasks=160527
Total megabyte-milliseconds taken by all map tasks=164379648
Map-Reduce Framework
Map input records=1
Map output records=1
Input split bytes=546
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=10322
CPU time spent (ms)=7470
Physical memory (bytes) snapshot=487501824
Virtual memory (bytes) snapshot=8262631424
Total committed heap usage (bytes)=276299776
File Input Format Counters
Bytes Read=0
File Output Format Counters
Bytes Written=0
18/06/06 00:45:41 INFO mapreduce.ExportJobBase: Transferred 660 bytes in 58.8163 seconds (11.2214 bytes/sec)
18/06/06 00:45:41 INFO mapreduce.ExportJobBase: Exported 1 records.
[acadmild@192 ~]$
```

Sqoop Assignment

Select * from person; database SqoopExport2

```
mysql> create database SqoopExport2;  
Query OK, 1 row affected (0.04 sec)
```

```
mysql> use SqoopExport2;  
Database changed
```

```
mysql> Create table Person(  
-> person_id INT NOT NULL,  
-> lname varchar(20),  
-> fname varchar(20),  
-> area varchar(20),  
-> city varchar(20),  
-> PRIMARY KEY (person_id)  
-> );  
Query OK, 0 rows affected (0.23 sec)
```

```
mysql> select * from Person;  
+-----+-----+-----+-----+-----+  
| person_id | lname | fname | area | city |  
+-----+-----+-----+-----+-----+  
| 3 | James | Brown | New York | United States |  
+-----+-----+-----+-----+-----+  
1 row in set (0.00 sec)
```

```
mysql>
```

acad... [sbin... [had... mys... [ACD... ACD... *ne... acad... acad...